

# STA2201H Methods of Applied Statistics II

Monica Alexander

Week 10: Splines

# Notes

- ▶ Last lab today, but you can drop the lowest grade
- ▶ Outstanding lab grades back by tomorrow
- ▶ A2 grades back by end week

## Reading

- ▶ Eilers, P.H. and Marx, B.D., 1996. Flexible smoothing with B-splines and penalties. Statistical science, pp.89-102.
- ▶ Friedman, J., Hastie, T. and Tibshirani, R., 2001. The elements of statistical learning. Chapter 5
- ▶ James, G., Witten, D., Hastie, T. and Tibshirani, R., 2013. An introduction to statistical learning. Chapter 7
- ▶ BDA Chapter 20

# The end goal for today

What are we doing: Bayesian Penalized B-Splines (P-Splines) Regression

How do we get there

1. What are splines?
2. What are B-splines?
3. How to fit B-Splines regression?
4. How are they penalized?
5. How to fit P-Splines in a Bayesian framework?

## Moving beyond linearity

- ▶ Replace the vector of inputs  $X$  with transformations of  $X$
- ▶ Use linear models in the new space of derived inputs

i.e. instead of fitting a linear model in  $X$  we fit the model

$$y_i = \alpha_0 + \alpha_1 b_1(x_i) + \alpha_2 b_2(x_i) + \cdots + \alpha_K b_K(x_i) + \epsilon_i$$

- ▶ the functions  $b_k(.)$  are called **basis functions** and are fixed and known i.e. we choose them.
- ▶ can think of this as a standard linear model so all the usual techniques apply (least squares, standard errors on coefficients, etc)

# Basis functions

We have

$$y_i = \mu(x_i) + \varepsilon_i$$

$$\varepsilon_i \sim N(0, \sigma_y^2)$$

with  $\mu(x_i) = \sum_{k=1}^K b_k(x_i)\alpha_k$ .

What to choose for  $b(\cdot)$ ?

- ▶ simplest: piece-wise constant
- ▶ piece wise polynomials
- ▶ ... regression splines

# Piecewise polynomials


- ▶ Divide the domain of  $X$  into  $K$  intervals defined by **knot points**  $\tau_1, \tau_2, \dots, \tau_{K+1}$ , and represent  $y$  as a separate polynomial in each interval

E.g. piecewise constant with two knot points

$$b_1(X) = I(X < \tau_1), \quad b_2(X) = I(\tau_1 \leq X < \tau_2), \quad b_3(X) = I(\tau_2 \leq X)$$

## HTF fig 5.1

### Piecewise Constant






# Piecewise polynomials

E.g. piecewise linear

$$b_{m+3} = b_m(X)X, m = 1, \dots, 3$$

Piecewise Linear



(HTF Fig 5.1)

## Piecewise polynomials with constraints

- ▶ We would usually want to constrain the pieces to be continuous at the knot points

e.g. in the piecewise linear case

$$b_1(X) = 1, \quad b_2(X) = X, \quad b_3(X) = (X - \tau_1)_+, \quad b_4(X) = (X - \tau_2)_+$$

$$\text{where } (X - \tau_1)_+ = \begin{cases} (X - \tau_1) & \text{if } x > \tau_1 \\ 0 & \text{otherwise} \end{cases}$$

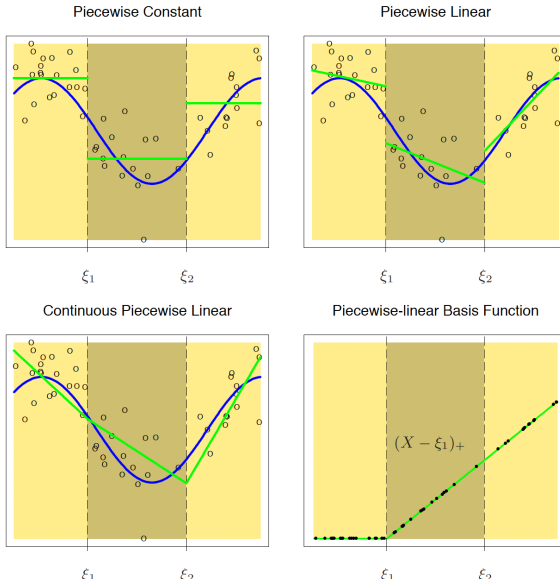
Remember that these add so

$$y_i = \sum_K \alpha_k b_k(x_i) + \epsilon_i$$

So the model is linear with changing slope.

## HTF fig 5.1

Bottom right hand panel is the function  $(X - \tau_1)_+$



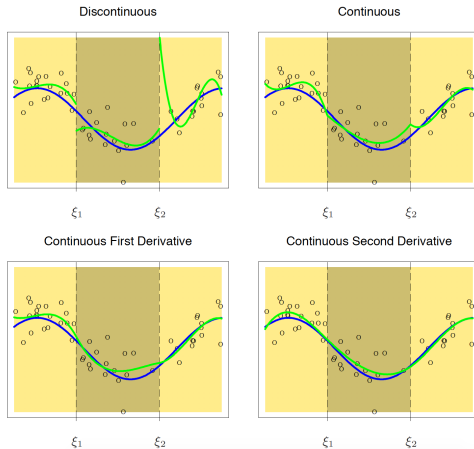
# Regression splines

- ▶ Piece-wise polynomials with continuity restrictions are called splines.
- ▶ In particular, a degree- $d$  spline is a piecewise degree- $d$  polynomial with continuity in derivatives up to degree  $d - 1$  at each knot.
- ▶ We often prefer smoother functions; these can be achieved by increasing the order of the local polynomial

Commonly chosen are **cubic splines**: third-order polynomial with continuity in first and second derivatives

# Cubic polynomials

HTF, Figure 5.2. Bottom RH panel is a cubic spline



## Cubic spline basis

For the case with two knots,

$$\begin{aligned} b_1(X) &= 1, & b_3(X) &= X^2, & b_5(X) &= (X - \tau_1)_+^3 \\ b_2(X) &= X, & b_4(X) &= X^3, & b_6(X) &= (X - \tau_2)_+^3 \end{aligned}$$

Easy to show this is indeed a representation of a cubic spline (write out, show that in each knot, the relevant pair of functions has identical function values, and identical first and second derivatives.)

## Cubic splines

“It is claimed that cubic splines are the lowest-order spline for which the knot-discontinuity is not visible to the human eye. There is seldom any good reason to go beyond cubic-splines, unless one is interested in smooth derivatives.” HTF page 144

## Back to the check list

What are we doing: Bayesian Penalized B-Splines Regression

1. What are splines?
2. What are B-splines?
3. How to fit B-Splines regression?
4. How are they penalized?
5. How to fit P-Splines in a Bayesian framework?



# An alternative basis representation

- ▶ The previous representation of splines is called the **truncated power basis**
- ▶ There are many equivalent bases for representing splines
- ▶ The truncated power basis is conceptually simple, but not great numerically (we want to avoid potentially taking powers of large numbers)

Alternative: **B-spline basis**

## B-spline basis

- ▶ Define a set of knot points  $\tau_1 < \tau_2 < \dots < \tau_{K+1}$
- ▶ Denote by  $B_{i,m}(x)$  the  $i$ th B-spline basis function of degree  $m$  for the knot sequence  $\tau$
- ▶ The B-splines are defined recursively

For degree 0:

$$B_{i,0}(x) = \begin{cases} 1 & \text{if } \tau_i \leq x < \tau_{i+1} \\ 0 & \text{otherwise} \end{cases}$$

Then

$$B_{i,m}(x) = \frac{x - \tau_i}{\tau_{i+m} - \tau_i} B_{i,m-1}(x) + \frac{\tau_{i+m+1} - x}{\tau_{i+m+1} - \tau_{i+1}} B_{i+1,m-1}(x)$$

## B-splines

$$B_{i,0}(x) = \begin{cases} 1 & \text{if } \tau_i \leq x < \tau_{i+1} \\ 0 & \text{otherwise} \end{cases}$$

Then


$$B_{i,m}(x) = \frac{x - \tau_i}{\tau_{i+m} - \tau_i} B_{i,m-1}(x) + \frac{\tau_{i+m+1} - x}{\tau_{i+m+1} - \tau_{i+1}} B_{i+1,m-1}(x)$$

- ▶  $B_{i,0}(x)$  is piecewise constant, indicating which knot span  $x$  is in
- ▶ In the recursion formula, the first part ramps up as  $x$  goes from  $\tau_i$  to  $\tau_{i+m}$
- ▶ The second part ramps down as  $x$  goes from  $\tau_{i+1}$  to  $\tau_{i+m+1}$

# Linear B-splines

- For example,  $B_{i,1}(x)$  is a triangular function that is zero below  $x = \tau_i$ , ramps to one at  $x = \tau_{i+1}$  and back to zero at and beyond  $x = \tau_{i+2}$ .


Linear splines  
knots shown with red dotted lines



# Linear B-splines


For linear splines

- ▶ Splines consists of 2 pieces, each of degree 1
- ▶ the pieces join at 1 knot
- ▶ Splines are non-zero on a domain spanning 3 knots
- ▶ Except at boundaries, overlap with 2 of their neighbors



# Cubic B-splines

Cubic splines  
knots shown with red dotted lines



# B-Splines


In general, a  $d$ -degree B-Spline has the general properties

- ▶ consists of  $d + 1$  polynomial pieces, each of degree  $d$
- ▶ the polynomial pieces join at  $d$  inner knots
- ▶ at the joining points, derivatives up to order  $d - 1$  are continuous
- ▶ the B-spline is positive on a domain spanned by  $d + 2$  knots; everywhere else it is zero;
- ▶ except at the boundaries, it overlaps with  $2d$  polynomial pieces of its neighbors;
- ▶ at a given  $x$ ,  $d + 1$  B-splines are nonzero
- ▶ at any given  $x$ ,  $\sum_{k=1}^K B_k(x) = 1$

# Generating B-Splines in R

Can just use bs function as part of the splines package as a starting point. e.g. the linear splines plot

```
library(splines)
x <- seq(0,10, by = 0.1)
knots <- seq(0,10, by = 2)
deg1 <- bs(x, degree = 1, knots = knots)
colnames(deg1) <- paste0("b", 1:(length(knots)+1))
deg1 <- as_tibble(deg1)
deg1 %>%
  mutate(x = x) %>%
  pivot_longer(-x, names_to = "spline", values_to = "value") %>%
  ggplot(aes(x, value, color = spline)) +
  geom_line(lwd = 1.5) +
  geom_vline(xintercept = knots, color = "red", lty = 2, lwd = 1.2)+
  theme_bw(base_size = 20) +
  ggtitle("Linear splines\nknots shown with red dotted lines")
```





## Back to the check list

What are we doing: Bayesian Penalized B-Splines Regression

1. What are splines?
2. What are B-splines?
3. How to fit B-Splines regression?
4. How are they penalized?
5. How to fit P-Splines in a Bayesian framework?

## How does this help us?

Model outcome  $y$  as

$$y_i = \mu(x_i) + \varepsilon_i$$


$$\varepsilon_i \sim N(0, \sigma_y^2)$$

with  $\mu(x_i) = \sum_{k=1}^K B_k(x_i)\alpha_k$ .

- ▶ choose knot points  $k$ , degree of splines
- ▶ generate basis splines  $B_k$  based on  $x_i$ 's, fit regression to get estimates of  $\alpha_k$ 's
- ▶ Different  $\alpha_k$ 's give different  $\mu(x)$ s

## Example (Lesaffre and Lawson (2012) Chapter 10)

- ▶ Modeling the relationship between osteoporosis and BMI
- ▶ Osteoporosis is a disease where the bone mineral density is reduced so that the bones become fragile.
- ▶ Marker used is total body bone mineral content (TBBMC, in kg).



# Fit B-Splines regression

fit the model

$$\begin{aligned}y_i &= \mu(x_i) + \varepsilon_i \\ \mu(x_i) &= \sum_{k=1}^K B_k(x_i) \alpha_k \\ \varepsilon_i &\sim N(0, \sigma_y^2)\end{aligned}$$

where  $y$  is TBBMC and  $x$  is BMI. Let's use equally-spaced knots every 2.5 increments of BMI

# Get B-Splines

```
I <- 2.5 # between-knot length  
res <- getsplines(x.i, I = I) # a function from distortr, to get splines of constant shape  
B.ik <- res$B.ik  
x.i[1] # look at first value of BMI
```

```
## [1] 23.61275
```

```
round(B.ik[1,],2) # value of splines at x.i[1]
```

```
## spline1 spline2 spline3 spline4 spline5 spline6 spline7 spline8  
## 0.00 0.00 0.00 0.16 0.67 0.17 0.00 0.00  
## spline9 spline10 spline11 spline12 spline13  
## 0.00 0.00 0.00 0.00 0.00
```

```
sum(B.ik[1,]) # confirm = 1
```


```
## [1] 1
```

```
round(res$knots.k,1) # indicates where spline is at max
```

```
## spline1 spline2 spline3 spline4 spline5 spline6 spline7 spline8  
## 13.6 16.1 18.6 21.1 23.6 26.1 28.6 31.1  
## spline9 spline10 spline11 spline12 spline13  
## 33.6 36.1 38.6 41.1 43.6
```

# Get B-Splines

red dot is  $x_1$




# Fit in Stan

```
data {  
  int<lower=0> N;  
  int<lower=0> K;  
  vector[N] y;  
  matrix[N,K] B;  
}  
parameters {  
  vector[K] alpha;  
  real<lower=0> sigma;  
}  
transformed parameters{  
  vector[N] mu;  
  mu = B*alpha;  
}  
model {  
  
  //likelihood  
  y ~ normal(mu, sigma);  
  //priors  
  alpha ~ normal(0,1);  
  sigma ~ normal(0,1);  
}
```

# Results

Note, a good example of where `gather_draws` and `median_qi` is quick and easy

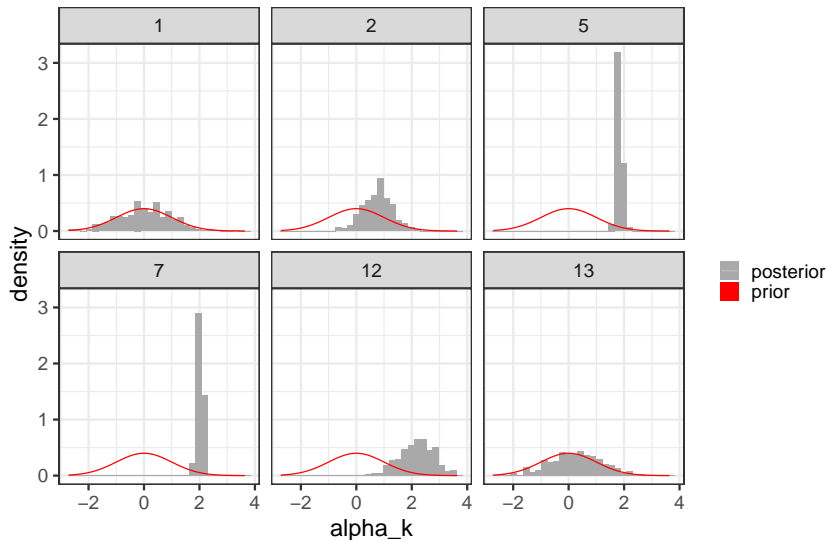
```
mod1 %>%  
  gather_draws(mu[condition]) %>%  
  median_qi() %>%  
  ungroup() %>%  
  mutate(tbbmc = y.i, bmi = x.i) %>%  
  ggplot(aes(bmi, tbbmc)) + geom_point(aes(color = "data")) +  
  geom_point(aes(bmi, .value, color = "fit")) +  
  geom_errorbar(aes(ymin = .lower, ymax = .upper, color = "fit"), alpha = 0.2) +  
  scale_color_manual(name = "", values = c("fit" = "red", "data" = "black")) + theme_bw(base_size = 20)
```





The closer to the boundary, the more the  $\alpha_k$  are informed by prior

Prior and posterior densities of selected alphas



## Obtaining estimates for all $x$

- ▶ Can we estimate  $\mu(x)$  for some  $x$  that was not in the data set, e.g.  $x = 36$
- ▶ Yes, we can evaluate  $\mu(x)$  at a grid of  $x$  values (as long as they are in the range of observed  $x$  values (such that  $\alpha_k$ 's have been estimated)).
- ▶ Get posterior samples using posterior samples for  $\alpha_k$ 's

$$\mu(x)^{(s)} = \sum_{k=1}^K B_k(x) \alpha_k^{(s)}$$


As usual, could do this in Stan or in R.

## Obtaining estimates for “all” x

- Need to get values of B-Splines at each grid point then we are good to go

```
res_full <- getsplines(seq(min(x.i),max(x.i), by = 0.1), I = I)
B.ik_full <- res_full$B.ik #splines at full grid
alpha.sk <- extract(mod1)["alpha"] #post samples of alpha
CI.iq <- matrix(NA, nrow(B.ik_full),3)
for (i in 1:nrow(CI.iq)){
  CI.iq[i,] <- quantile(B.ik_full[i,]%*%t(alpha.sk),
                        c(0.025, 0.5, 0.975))
}
```

## Obtaining estimates for “all” $x$



Note: what's the difference between generating new  $\mu(x)$  versus new  $y$

## Back to the check list

What are we doing: Bayesian Penalized B-Splines Regression


1. What are splines?
2. What are B-splines?
3. How to fit B-Splines regression?
4. How are they penalized?
5. How to fit P-Splines in a Bayesian framework?

# Choices


1. What degree splines should we use to fit (i.e. what should  $d$  be)?
  - ▶ in general, as mentioned above, should be at most cubic ( $d = 3$ )
2. What to do at boundaries
  - ▶ Above I used the same splines at the boundaries, resulting in some splines which overlap with the data range only partially.
  - ▶ This is easier to interpret when we get into penalization
  - ▶ Alternative implementations include fewer splines with higher values at the boundaries.

# Choices

Cubic splines  
same shape splines at boundaries



Cubic splines  
different splines at boundaries



## Choices

### 3. Where to put the knots?

- ▶ Options include equally spaced or based on data availability (e.g., based on percentiles of data set)
- ▶ I prefer equally spaced splines
  - ▶ to avoid limited uncertainty for  $x$ -values where data are sparse
  - ▶ to facilitate interpretation when a penalization is used
  - ▶ ... but as usual, sensitivity to model choice (in this case, knot placement) should be checked

### 4. How many knots?

- ▶ The regression is most flexible in regions that contain a lot of knots, because in those regions the coefficients can change rapidly
- ▶ Distance needs to be small enough to capture (true) fluctuations in outcome of interest ( $\mu(x)$ )
- ▶ But too many knots may be just picking up random fluctuations



## Smoothing penalties

- ▶ In fitting a smooth curve to a set of data, we want the residual sum of squares to be small

$$\text{RSS}(f) = \sum_{i=1}^n \{y_i - f(x_i)\}^2$$

where  $f(x_i) = \sum_{k=1}^K \alpha_k B_k(x_i)$ .

- ▶ But if we don't put any constraints on  $\sum_{k=1}^K \alpha_k B_k(x_i)$ , we can make  $\text{RSS}$  zero by just interpolating all the points
- ▶ If the number of knots be relatively large, such that the fitted curve will show more variation than is justified by the data
- ▶ What we really want is a function that makes  $S$  small but is also **smooth**
- ▶ Add another term that penalizes fluctuations in spline

## Smoothing splines

- ▶ Penalized RSS (O'Sullivan, Hastie)

$$\text{RSS}(f, \lambda) = \sum_{i=1}^N \{y_i - f(x_i)\}^2 + \lambda \int \{f''(t)\}^2 dt$$

where  $\lambda$  is a fixed smoothing parameter. Loss+penalty / bias v variance trade-off:

- ▶ The smaller the  $\lambda$ , the closer  $f$  is to interpolation
- ▶ As  $\lambda \rightarrow \infty$ , we get a simple least squares line fit.
- ▶  $\int \{f''(t)\}^2 dt$  a measure of total change in the derivative
- ▶ The bigger  $\int \{f''(t)\}^2 dt$ , the more 'rough' the fit
- ▶ It can be shown that the function that minimizes this penalized RSS is a natural cubic spline (cubic spline + linear in the region outside the extreme knots) with knots at each  $x_i$ . Then choose  $\lambda$  based on LOO-CV. (HTF Chap 5)

## Alternative: P-Splines

Eilers and Marx (1996) suggested an alternative smoothing penalty:

$$\text{RSS}(f, \lambda) = \sum_{i=1}^n \{y_i - f(x_i)\}^2 + \lambda \sum_{k=q+1}^K (\Delta^q \alpha_k)^2$$

- ▶ the penalty is based on finite differences of the coefficients of adjacent B-splines
- ▶ This reduces the dimensionality of the problem to  $K$ , the number of B-splines, instead of  $n$ , the number of observations
- ▶ A good discrete approximation to term used in the previous slide
- ▶  $q$  could be anything, in practice usually consider penalizing first- or second-order differences
- ▶ These are called penalized B-splines, or P-splines

## P-splines

In likelihood-based inference, want to maximize the penalized likelihood

$$L = l(y, \alpha_1, \dots, \alpha_k) - \lambda \sum_{k=q+1}^K (\Delta^q \alpha_k)^2$$

This leads to the system of equations

$$B'(y - \mu) = \lambda D_q' D_q \alpha$$

- ▶  $D_q$  is the matrix representation of the difference operator, e.g.  $D_1$  is dimensions  $(k-1) \times k$  and has elements  $d_{ij} = -1$  if  $i = j$ ,  $d_{ij} = 1$  if  $i = j - 1$  and 0 otherwise
- ▶ Elements of  $B$  are  $B_k(x_i)$
- ▶ Note if  $\lambda = 0$ , just usual linear regression with B-spline basis. If  $q = 0$ , this is ridge regression
- ▶ For fixed  $D_q$  and  $\lambda$ , can solve using IWLS (lecture 2)
- ▶ Choose smoothing parameter  $\lambda$  based on AIC or CV

## Back to the check list

What are we doing: Bayesian Penalized B-Splines Regression

1. What are splines?
2. What are B-splines?
3. How to fit B-Splines regression?
4. How are they penalized?
5. How to fit P-Splines in a Bayesian framework?

# Bayesian P-Splines

Our penalized likelihood is

$$L = l(y, \alpha_1, \dots, \alpha_k) - \lambda \sum_{k=q+1}^K (\Delta^q \alpha_k)^2$$

- ▶ In the Bayesian context, the spline coefficients  $\alpha_k$  are now random variables, which need prior distributions
- ▶ The stochastic equivalent to the difference penalties is to place a q-order random walk prior on the  $\alpha_k$ s

# Bayesian P-Splines

- ▶ E.g. for a first-order difference penalty, the prior is

$$\alpha_k = \alpha_{k-1} + \varepsilon_k$$
$$\varepsilon_k | \sigma \sim N(0, \sigma_\alpha^2)$$

or equivalently

$$\alpha_k \sim N(\alpha_{k-1}, \sigma_\alpha^2)$$

(We saw these last week in the context of temporal models!)


- ▶ The  $\sigma_\alpha^2$  term is equivalent to  $\lambda$ , and controls the smoothness of fit
- ▶ The smaller  $\sigma_\alpha^2$ , the smoother the fit

# How to fit in Stan


```
data {  
  int<lower=0> N;  
  int<lower=0> K;  
  vector[N] y;  
  matrix[N,K] B;  
}  
parameters {  
  vector[K] alpha;  
  real<lower=0> sigma;  
  real<lower=0> sigma_alpha;  
}  
transformed parameters{  
  vector[N] mu;  
  mu = B*alpha;  
}  
model {  
  //likelihood  
  y ~ normal(mu, sigma);  
  //priors  
  alpha[1] ~ normal(0, sigma_alpha);  
  alpha[2:K] ~ normal(alpha[1:(K - 1)], sigma_alpha);  
  alpha ~ normal(0,1);  
  sigma ~ normal(0,1);  
  sigma_alpha ~ normal(0,1);  
}
```



## Example



But what if we had a lot less data



... and second-order penalization would be even more smooth.

## Post-script: GAMs

- ▶ In the example above we were using (P)-splines to estimate the relationship between outcome of interest ( $y$ , in this case TBBMC) and one covariate  $x$
- ▶ Could extend to model to include additional covariates, and potentially model in the same way (or in the usual 'fixed effects' way)
- ▶ And in doing so, we have arrived at an example of a generalized additive model

# GAMs

- ▶ Assume that given a set of covariates  $x_i$ ,  $y_i$  has a distribution that belongs to the exponential family.
- ▶ We have a link function  $g$  such that  $\mu_i = E(y_i|x_i) = g^{-1}(\eta_i)$
- ▶ Recall that in the GLM case the  $\eta_i$  was our linear predictor
- ▶ In the GAMs case,  $\eta_i$  is a additive predictor

$$\eta_i = f_1(x_{i1}) + \cdots + f_p(x_{ip})$$


- ▶ the functions  $f_p$  are unknown smooth functions of the covariates
- ▶ Some covariates may have known functions, in which case our model is semi-parametric
- ▶ We may have up to  $p$  smoothing parameters  $\lambda$ 's

Longer post-script: P-splines for temporal  
smoothing

# Temporal smoothing

- ▶ Last week we talked about models to fit and project time series
- ▶ E.g. we fit a (hierarchical) RW2 model to estimate and project foster care populations in the US

Estimated and projected entries per capita  
hierarchical second-order random walk



# Temporal smoothing

- ▶ We could have used P-splines in same way
- ▶ Fit a model with second-order P-splines

$$y_{st} \sim N(\log \lambda_{st}, \sigma_y^2)$$

$$\log \lambda_{st} = \alpha_k B_k(t)$$

with

$$\Delta^2 \alpha_k \sim N(0, \sigma_{\alpha,s}^2)$$

- ▶  $y_{st} = \log(E_{st}/P_{st})$ ,  $E$  is entries,  $P$  is population
- ▶ For comparison, the model we fit last week was essentially  $\Delta^2 \log \lambda_{st} \sim N(0, \sigma_\lambda^2)$ .
- ▶ RW2 process now on coefficients, not data
- ▶ Projection happens through projection of the coefficients

# Temporal smoothing

$$y_{st} \sim N(\log \lambda_{st}, \sigma_y^2)$$

$$\log \lambda_{st} = \alpha_k B_k(t)$$

with

$$\Delta^2 \alpha_k \sim N(0, \sigma_{\alpha,s}^2)$$

In addition, want to model variance hierarchically

$$\log \sigma_{\alpha,s} \sim N(\mu_\sigma, \tau^2)$$

- ▶ with P-splines, usually choose to have a relatively large number of splines (knot points,  $k$ ), then smooth away fluctuations
- ▶ I ran model with knots every 2.5 years, with constant spacing across all states and boundary splines that are same shape
- ▶ Helps to interpret  $\sigma_\alpha$ 's as smoothing parameters, given the same spline settings



# How to get P-spline projections?

In R to spell it out (could do in Stan)

```
proj_years <- 2018:2030

# Note: B.ik are splines for in-sample period
# has dimensions i (number of years) x k (number of knots)

# need splines for whole period
B.ik_full <- getsplines(c(years, proj_years))$B.ik

K <- ncol(B.ik) # number of knots in sample
K_full <- ncol(B.ik_full) # number of knots over entire period
proj_steps <- K_full - K # number of projection steps

# get your posterior samples
alphas <- extract(mod)[["alpha"]]
sigmas <- extract(mod)[["sigma"]] # sigma_alpha
sigma_ys <- extract(mod)[["sigma_y"]]

nsims <- nrow(alphas)
```

```

# first, project the alphas
alphas_proj <- array(NA, c(nsim, proj_steps, length(states)))
set.seed(1098)
# project the alphas
for(j in 1:length(states)){
  first_next_alpha <- rnorm(n = nsims,
    mean = 2*alphas[,K,j] - alphas[,K-1,j],
    sd = sigmas[,j])
  second_next_alpha <- rnorm(n = nsims,
    mean = 2*first_next_alpha - alphas[,K,j],
    sd = sigmas[,j])

  alphas_proj[,1,j] <- first_next_alpha
  alphas_proj[,2,j] <- second_next_alpha
  # now project the rest
  for(i in 3:proj_steps){ #!!! not over years but over knots
    alphas_proj[,i,j] <- rnorm(n = nsims,
      mean = 2*alphas_proj[,i-1,j] - alphas_proj[,i-2,j],
      sd = sigmas[,j])
  }
}


# now use these to get y's
y_proj <- array(NA, c(nsim, length(proj_years), length(states)))
for(i in 1:length(proj_years)){ # now over years
  for(j in 1:length(states)){
    all_alphas <- cbind(alphas[,.j], alphas_proj[,.j] )
    this_lambda <- all_alphas %*% as.matrix(B.ik_full[length(years)+i, ])
    y_proj[,i,j] <- rnorm(n = nsims, mean = this_lambda, sd = sigma_ys[,j])
  }
}

# then proceed as normal to get median, quantiles etc

```

## P-splines projection

Estimated and projected entries per capita  
second-order P-splines



## Temporal take-aways

- ▶ From a time series perspective, P-splines are a nice way of capturing non-linear trends over time
- ▶ Temporal structure is on coefficients (from knot point to knot point) rather than on data (from time point to time point)
- ▶ Can hierarchically smooth fluctuations in the same way we did with temporal models last week
- ▶ Particularly powerful in combination with covariates (outcome = expected + smoothed fluctuations)
- ▶ Compared to random walks or ARIMA processes, will have different uncertainty in fit and projections
- ▶ Not much different in foster care case, but becomes more obvious when data are sparse